

# OEAW AI SUMMER SCHOOL

## DEEP LEARNING II

### Neural Networks



Johannes Brandstetter, Michael Widrich  
Institute for Machine Learning

# Lecture II: Neural Networks



**Perceptron**

**Multi-Layer Perceptron**

# Neurophysiological background

- The inside of every neuron (nerve or brain cell) carries a certain **electric charge**.
- Electric charges of connected neurons may raise or lower this charge:
  - by means of transmission of ions through the synaptic interface.
- As soon as the charge reaches a certain **threshold**, an **electric impulse is transmitted** through the cell's axon to the neighboring cells.
- In the synaptic interfaces, chemicals called neurotransmitters **control the strength to which an impulse is transmitted** from one cell to another.

# Perceptrons

- A perceptron is a simple linear threshold unit:


$$g(\mathbf{x}_i; \mathbf{W}, \theta) = \begin{cases} 1 & \text{if } \mathbf{W}^T \mathbf{x}_i > \theta \\ 0 & \text{otherwise} \end{cases}$$

- In analogy to the biological model:
  - **inputs**  $\mathbf{x}_i \Rightarrow$  charges received from connected cells
  - **weights**  $\mathbf{W} \Rightarrow$  properties of the synaptic interface
  - **output**  $\Rightarrow$  impulse that is sent through the axon as soon as the charge exceeds the threshold  $\theta$
- Though it seems to be a (simplistic) model of a neuron, a **perceptron is nothing else but a simple linear classifier.**

# Perceptrons and linear separability

- In case that the data set  $Z$  is **linearly separable**, the perceptron learning algorithm terminates and finally solves the learning task.
- The final solution is **not unique**:
  - arbitrary solution (depending on initial weights)
- Perceptrons cannot solve classification tasks that are not linearly separable:
  - simple XOR problems
- Solution of introducing **intermediate layers**:
  - The outputs of the first layer are used as **input of the first intermediate layer**.

# Multi-layer perceptrons



## Some historical remarks


- Minsky and Papert in the late 1960s: a training algorithm for MLPs is computationally infeasible.
  - Study of multi-layer perceptrons is not worthwhile.
  - The study of multi-layer perceptrons was almost halted until the mid of the 1980s.
- In 1986, Rumelhart, McClelland, Hinton first published the backpropagation algorithm.
  - Similarly described by Werbos (1974) and Bryson (1960s).
  - The key idea is to replace the discontinuous threshold function by a differentiable function. Then the output of the neuron, its so-called activation, is computed as:

$$g(\mathbf{x}; \mathbf{W}, \theta) = \varphi(\mathbf{W}^T \mathbf{x} - \theta)$$




# Activation functions

- MLP is **fully connected feed-forward neural network (FNN)**:
  - Today, “FNN” and “MLP” are often used synonymously.
  - Convolutional NNs are only partially connected, Recurrent NNs have feedback loops (No MLPs!).
- Without activation functions, neural networks are linear:
  - **Activations allow to learn nonlinear functions.**
- Earlier used activation functions: sigmoid, tanh
  - Sigmoid function  $\sigma(x)$
  - $\tanh(x) = 2\sigma(2x) - 1$
  - Rescaled versions of each other



# Learning non-linear functions



$x_1$	$x_2$	$y$
0	0	0
0	1	1
1	0	1
1	1	0

This classification problem:

- ... can't be solved by logistic regression.
- ... is trivial to solve using an MLP.

# How powerful are neural networks?

- Useful measure of a machine learning algorithm: what are the **most complicated functions it can learn**?

- Measured by the **VC dimension**<sup>1</sup>

- VC dimension for a given neural network:

$$d_{VC} \leq O(W \log(c \cdot M))$$

(W = number of weights, M = number of units)

- VC dimensions of neural networks are much higher than for e.g. Logistic Regression.

- NNs can learn (much) more complex decision functions.

- **Neural Networks are “Universal Function Approximators”.**

---

<sup>1</sup>**Vapnik-Chervonenkis** dimension – to learn more, see “Statistical Learning Theory” (Vladimir N. Vapnik)

# Universal Approximation Theorem

*For any given continuous function  $f \in I_m \equiv [0, 1]^m$  and  $\varepsilon > 0$ , there exists a function of the form*

$$F(x) = \sum_{i=1}^N \alpha_i \varphi(w_i^T x + b_i)$$

*such that*

$$|F(x) - f(x)| < \varepsilon$$

*for all  $x \in I_m$ .*


For the proof, see [Hornik (1991) "Approximation Capabilities of Multilayer Feedforward Networks"]

# **Backpropagation**

## **Vanishing Gradient**


## Remember: Logistic Regression

$$g(\mathbf{x}_i) = \sigma(\mathbf{W}^T \mathbf{x}_i + b)$$



Logistic regression with **learned features**:

$$g(\mathbf{x}_i) = \sigma(\mathbf{W}_{(2)}^T \mathbf{f}_i + b)$$




# Gradients for Logistic Regression

- **Using:**  $\frac{\partial \sigma(z)}{\partial z} = \sigma(z) \cdot (1 - \sigma(z))$
- **Using**  $\sigma_i$  instead of  $\sigma(\mathbf{W}^T \mathbf{x}_i)$

$$\begin{aligned} L &= - \sum_i \left( y_i \log \sigma(\mathbf{W}^T \mathbf{x}_i) + (1 - y_i) \log(1 - \sigma(\mathbf{W}^T \mathbf{x}_i)) \right) \\ \frac{\partial L}{\partial \mathbf{W}} &= - \sum_i \left( y_i \frac{1}{\sigma_i} \cdot \sigma_i \cdot (1 - \sigma_i) \cdot \mathbf{x}_i - \frac{1 - y_i}{1 - \sigma_i} \cdot \sigma_i \cdot (1 - \sigma_i) \cdot \mathbf{x}_i \right) \\ &= - \sum_i \left( y_i(1 - \sigma_i) \cdot \mathbf{x}_i - (1 - y_i) \cdot \sigma_i \cdot \mathbf{x}_i \right) \\ &= - \sum_i (y_i - y_i \sigma_i - \sigma_i + \sigma_i y_i) \mathbf{x}_i \\ &= \sum_i (\sigma_i - y_i) \mathbf{x}_i \end{aligned}$$

# Neural networks are nested structures


- Only **fully-connected feed-forward** connections:



$$y = f\left(h_5\left(h_4\left(h_3\left(h_2\left(h_1\left(x; \mathbf{W}_1\right); \mathbf{W}_2\right); \mathbf{W}_3\right); \mathbf{W}_4\right); \mathbf{W}_5\right); \mathbf{W}_6\right)$$




## One word how the matrices look like?



# Backpropagation of errors

- Loss errors are **backpropagated** to update the network:



$$y = f\left(h_5\left(h_4\left(h_3\left(h_2\left(h_1\left(\mathbf{x}; \mathbf{W}_1\right); \mathbf{W}_2\right); \mathbf{W}_3\right); \mathbf{W}_4\right); \mathbf{W}_5\right); \mathbf{W}_6\right)$$

# Backpropagation of errors

- Gradients multiply according to the [chain rule](#).
- Gradient signal gets lost in the noise of the network:
  - **Vanishing gradient**

$$\mathbf{W}_6 \leftarrow \mathbf{W}_6 - \eta \frac{\partial L}{\partial \mathbf{W}_6} \quad \eta \dots \text{learning rate}$$

$$\mathbf{W}_5 \leftarrow \mathbf{W}_5 - \eta \frac{\partial L}{\partial h_5} \frac{\partial h_5}{\partial \mathbf{W}_5}$$

$$\mathbf{W}_4 \leftarrow \mathbf{W}_4 - \eta \frac{\partial L}{\partial h_5} \frac{\partial h_5}{\partial h_4} \frac{\partial h_4}{\partial \mathbf{W}_4}$$



$$\mathbf{W}_3 \leftarrow \mathbf{W}_3 - \eta \frac{\partial L}{\partial h_5} \frac{\partial h_5}{\partial h_4} \frac{\partial h_4}{\partial h_3} \frac{\partial h_3}{\partial \mathbf{W}_3}$$

$$\mathbf{W}_2 \leftarrow \mathbf{W}_2 - \eta \frac{\partial L}{\partial h_5} \frac{\partial h_5}{\partial h_4} \frac{\partial h_4}{\partial h_3} \frac{\partial h_3}{\partial h_2} \frac{\partial h_2}{\partial \mathbf{W}_2}$$

$$\mathbf{W}_1 \leftarrow \mathbf{W}_1 - \eta \frac{\partial L}{\partial h_5} \frac{\partial h_5}{\partial h_4} \frac{\partial h_4}{\partial h_3} \frac{\partial h_3}{\partial h_2} \frac{\partial h_2}{\partial h_1} \frac{\partial h_1}{\partial \mathbf{W}_1}$$


# Activation functions = the main culprits

- Gradients multiply according to the chain rule.
- In the **saturation regions** of the activation functions the gradients are close to zero.
- First identified 1991 by Sepp Hochreiter



# Vanishing Gradient

- Gradient signal gets lost in the noise:




$$y = f\left(h_5\left(h_4\left(h_3\left(h_2\left(h_1\left(\mathbf{x}; \mathbf{W}_1\right); \mathbf{W}_2\right); \mathbf{W}_3\right); \mathbf{W}_4\right); \mathbf{W}_5\right); \mathbf{W}_6\right)$$

# How to solve the vanishing gradient problem

## ■ Currently applied methods:

- **Gating** (mostly used in RNNs → see Lecture V)
- **Normalization**: distorts gradient, increases noise
  - Batch normalization, layer normalization, weight normalization
- **Activation functions** which avoid vanishing gradient
- Clever **weight initialization**
- Further **regularization** techniques

## ■ Clever usage is “Trick of the Trade” in Deep Learning



## Example: ReLUs and ELUs

### ■ Rectified linear unit, by Nair and Hinton in 2010:


- Idea:  $N$  sigmoids with shared weights but different biases:

$$\sum_i^N \sigma(h - i + 0.5) \approx \log(1 + e^h) \approx \max(0, h)$$

where  $h = \mathbf{w} \cdot \mathbf{x} + b$

- Gradient is either 0 or 1 (helps against vanishing gradients!)
- ReLU nets learn a piecewise linear function
- Problem: Dying ReLUs

### ■ Exponential linear unit, by Clevert, Unterthiner, Hochreiter in 2015:



# One word about loss functions

## ■ Regression:

- ☐ Mean-squared error for Gaussian noise assumption
- ☐ Absolute error for Laplace noise assumption
- ☐ If specific information about noise is known the loss function can be adapted.

## ■ Classification:

- ☐ Cross-entropy  $\Rightarrow$  maximizing the likelihood of correct classification
- ☐ Hinge loss (Support Vector Machines)



# **Deep Learning**

## **Trick of the Trade**

# Weight initialization

- 99 % of the time, a reasonably sized net will learn even with a cheap initialization.
- However:
  - Good initializations can help against **Vanishing Gradients**.
  - Good initializations can help to **converge quickly** (fewer iterations needed).
  - Good initializations are able to train even 30+ layer nets:
    - Not relevant in practice (way too deep!)
    - Interesting from research POV

# Weight initialization – simple schemes

- Simplest:  $\mathbf{W}_{ij} \sim \mathcal{N}(0, 0.01)$  (or some other small  $\sigma^2$ )
- LeCun 1998: Make sure layer output has variance of 1:

LeCun et al. Efficient backprop. Neural networks: Tricks of the trade.

- ☐ Depends on activation function
- ☐ Depends on number of input units  $k$
- ☐ Heuristic:  $\mathbf{W}_{ij} \sim \mathcal{U}(\frac{-1}{k}, \frac{1}{k})$

# Weight initialization – newer schemes

## ■ Glorot and Bengio 2010 (“Xavier initialization”)

Glorot, Bengio. Understanding the difficulty of training deep feedforward neural networks. AISTATS 2010

- Same variance between all layers in both forward pass (activations) and backward pass (delta errors).
- Boils down to  $k \cdot \text{Var}(\mathbf{W}) = \frac{1}{3}$
- Thus  $\mathbf{W}_{ij} \sim \mathcal{U}(\frac{-\sqrt{6}}{\sqrt{k+l}}, \frac{\sqrt{6}}{\sqrt{k+l}})$ :
  - $l$ : number of output units in a layer
  - $k$ : number of input units in a layer
  - Note: for ReLU, use  $\sqrt{3}$  instead of  $\sqrt{6}$

## ■ Saxe 2014 (“orthogonal initialization”)

Saxe et al. Exact solutions to the nonlinear dynamics of learning in deep linear neural networks. ICLR 2014

- Initialize with random orthogonal matrices
- Reasoning comes from analyzing linear nets
- Essential idea: keep determinant of Jacobian close to 1, then you won't lose much information

# Regularization in Deep Neural Networks


- **Overfitting** is an issue in neural networks.
  - Even more so in deep neural networks
- Good regularization schemes needed.
- Deep nets have enough units that each one can focus on one specific thing.
- Need to force units to prevent co-adaptation.

# Regularization in Deep Neural Networks


- **Overfitting** is an issue in neural networks.
  - Even more so in deep neural networks
- Good regularization schemes needed.
- Deep nets have enough units that each one can focus on one specific thing.
- Need to force units to prevent co-adaptation.
- **Presented regularization methods:**
  - Dropout
  - Batch normalization
  - Weight decay

# Dropout

- Idea: randomly “drop out” units during training
- Different units for each sample and in each iteration
- No unit can rely on the presence of other units for their work.
- Typical dropout rates: 0.5 for hidden and 0.2 for input units
- Implementation note: need to scale weights (or activations) after training



(a) Standard Neural Net



(b) After applying dropout.

Srivastava et al. Dropout: A Simple Way to Prevent Neural Networks from Overfitting. JMLR 2014

# Batch normalization

- During learning, the output distribution of a layer will change.
  - Higher layers have to continuously adapt to this change.
- BN normalizes each minibatch to mean 0 and std 1.
- On test time, use average mean/std of training set.
- BN helps both for regularization and optimization.
- Many variants exist:
  - Weight Norm
  - Layer Norm
  - ...
- Best way to do this is still not completely understood!



# Batch Normalization

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_1 \dots x_m\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.

**Input:** Network  $N$  with trainable parameters  $\Theta$ ;  
subset of activations  $\{x^{(k)}\}_{k=1}^K$

**Output:** Batch-normalized network for inference,  $N_{\text{BN}}^{\text{inf}}$

- 1:  $N_{\text{BN}}^{\text{tr}} \leftarrow N$  // Training BN network
- 2: **for**  $k = 1 \dots K$  **do**
- 3:   Add transformation  $y^{(k)} = \text{BN}_{\gamma^{(k)}, \beta^{(k)}}(x^{(k)})$  to  $N_{\text{BN}}^{\text{tr}}$  (Alg. 1)
- 4:   Modify each layer in  $N_{\text{BN}}^{\text{tr}}$  with input  $x^{(k)}$  to take  $y^{(k)}$  instead
- 5: **end for**
- 6: Train  $N_{\text{BN}}^{\text{tr}}$  to optimize the parameters  $\Theta \cup \{\gamma^{(k)}, \beta^{(k)}\}_{k=1}^K$
- 7:  $N_{\text{BN}}^{\text{inf}} \leftarrow N_{\text{BN}}^{\text{tr}}$  // Inference BN network with frozen parameters
- 8: **for**  $k = 1 \dots K$  **do**
- 9:   // For clarity,  $x \equiv x^{(k)}, \gamma \equiv \gamma^{(k)}, \mu_{\mathcal{B}} \equiv \mu_{\mathcal{B}}^{(k)}$ , etc.
- 10:   Process multiple training mini-batches  $\mathcal{B}$ , each of size  $m$ , and average over them:
$$\begin{aligned} \mathbb{E}[x] &\leftarrow \mathbb{E}_{\mathcal{B}}[\mu_{\mathcal{B}}] \\ \text{Var}[x] &\leftarrow \frac{m}{m-1} \mathbb{E}_{\mathcal{B}}[\sigma_{\mathcal{B}}^2] \end{aligned}$$
- 11:   In  $N_{\text{BN}}^{\text{inf}}$ , replace the transform  $y = \text{BN}_{\gamma, \beta}(x)$  with
$$y = \frac{\gamma}{\sqrt{\text{Var}[x] + \epsilon}} \cdot x + \left( \beta - \frac{\gamma \mathbb{E}[x]}{\sqrt{\text{Var}[x] + \epsilon}} \right)$$
- 12: **end for**

**Algorithm 2:** Training a Batch-Normalized Network

Ioffe, Szegedy. Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift.

# Weight decay

- Weights around 0 correspond to **low complexity models**.
- A higher model complexity necessitates higher weights (in terms of their absolute values).
  - Therefore, a mechanism that favors weights around 0 can be used to control model complexity.
  - We add  $\lambda \cdot \Omega(\mathcal{W})$  to the learning objective:
    - **Regularization term**  $\Omega(\mathcal{W})$  measures the overall size of the weights.
    - $\mathcal{W}$  is the set of all weights in the network.
    - **Regularization parameter**  $\lambda$  controls the influence of the regularization term.

## $L_2$ Weight Decay

- $\min_{\mathbf{W}} [\sum_{i=1}^n (g(\mathbf{x}_i; \mathbf{W}) - y_i)^2 + \lambda \sum_{j=1}^m w_j^2]$
- $\lambda = 0 \Rightarrow$  standard regression
- $\lambda = \infty \Rightarrow \mathbf{W} = 0$  (except  $w_0$ , straight line through mean)
- Equivalent: keep norm of  $\mathbf{W}$  smaller than some given constant:

$$\min_{\mathbf{W}} \sum_{i=1}^n (g(\mathbf{x}_i; \mathbf{W}) - y_i)^2$$
$$\text{s. t. } \sum_{j=1}^m w_j^2 \leq T$$

- This form of regularization has many names:
  - ☐  $L_2$  regularization (because it penalizes the  $L_2$  norm of  $\mathbf{W}$ )
  - ☐ Gaussian weight prior/decay
  - ☐ Ridge regression (only used for Regression, but not in e.g. neural networks)
  - ☐ Tikhonov regularization

# $L_1$ Weight Decay

- $\min_{\mathbf{W}} [\sum_{i=1}^n (g(\mathbf{x}_i; \mathbf{W}) - y_i)^2 + \lambda \sum_{j=1}^m |w_j|]$
- Penalizes  $L_1$  norm of  $\mathbf{W}$
- Has many names:
  - $L_1$  regularization
  - Laplace weight prior/decay
  - LASSO (least absolute shrinkage and selection operator)  
(only used for Regression, but not in e.g. neural networks)
  - Sparsity penalty term
- $L_2$  Regularization  $\rightarrow$  small parameters
- $L_1$  Regularization  $\rightarrow$  some parameters should be put exactly to 0
- “sparse” = “contains many zeros”

# Summary

- From perceptrons to multi-layer perceptron
  - Fully connected feed forward neural networks
- Learning non-linear functions
  - Activation functions
- Backpropagation of errors
  - Vanishing gradient
- Initialization
- Regularization:
  - Dropout
  - Batch normalization
  - Weight decay